



# Importing the Dependencies

```
In [ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from sklearn import metrics
from sklearn.linear_model import LinearRegression
from sklearn.svm import SVR
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error, roc_auc_score
```

```
In [ ]: import warnings
warnings.filterwarnings("ignore", message="X does not have valid feature names")
```

# Data Collection and Analysis

```
In [ ]: df = pd.read_csv('/content/ITC.NS.csv')
df.head()
```

```
Out[ ]:
```

|   | Date       | Open     | High     | Low      | Close    | Adj Close | Volume     |
|---|------------|----------|----------|----------|----------|-----------|------------|
| 0 | 01-01-1996 | 5.550000 | 5.600000 | 5.533333 | 5.583333 | 3.323907  | 985500.0   |
| 1 | 02-01-1996 | 5.466666 | 5.566666 | 5.288888 | 5.372222 | 3.198226  | 7470000.0  |
| 2 | 03-01-1996 | 5.133333 | 5.254444 | 5.101111 | 5.200000 | 3.095698  | 15160500.0 |
| 3 | 04-01-1996 | 5.200000 | 5.332222 | 5.144444 | 5.297777 | 3.153908  | 12397500.0 |
| 4 | 05-01-1996 | 5.297777 | 5.277777 | 5.188888 | 5.202222 | 3.097020  | 5008500.0  |

```
In [ ]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6890 entries, 0 to 6889
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   Date        6890 non-null  object  
1   Open        6879 non-null  float64  
2   High        6879 non-null  float64  
3   Low         6879 non-null  float64  
4   Close       6879 non-null  float64  
5   Adj Close   6879 non-null  float64  
6   Volume      6879 non-null  float64  
dtypes: float64(6), object(1)
memory usage: 376.9+ KB
```

```
In [ ]: df.shape
```

Out[ ]: (6890, 7)

```
In [ ]: df.describe()
```

```
Out[ ]:
```

|       | Open        | High        | Low         | Close       | Adj Close   | Volume  |
|-------|-------------|-------------|-------------|-------------|-------------|---------|
| count | 6879.000000 | 6879.000000 | 6879.000000 | 6879.000000 | 6879.000000 | 6.87900 |
| mean  | 121.511181  | 122.859461  | 120.011619  | 121.426830  | 99.025659   | 4.40242 |
| std   | 106.444634  | 107.405065  | 105.373130  | 106.383687  | 94.769531   | 9.92748 |
| min   | 4.182222    | 4.182222    | 4.144444    | 4.182222    | 2.489788    | 0.00000 |
| 25%   | 18.222221   | 18.541666   | 17.889999   | 18.169999   | 11.208004   | 8.77665 |
| 50%   | 72.599998   | 74.416664   | 71.000000   | 72.583336   | 51.009666   | 1.37196 |
| 75%   | 219.666672  | 222.266663  | 217.333328  | 219.933334  | 181.465515  | 2.58141 |
| max   | 432.799988  | 433.450012  | 429.350006  | 431.450012  | 431.450012  | 1.29416 |

## Transform Data

```
In [ ]: columns_to_2decimal = ['Open', 'High', 'Low', 'Close']
decimal_places = 2

df[columns_to_2decimal] = np.round(df[columns_to_2decimal], decimal_places)
df
```

```
Out[ ]:
```

|      | Date       | Open   | High   | Low    | Close  | Adj Close  | Volume     |
|------|------------|--------|--------|--------|--------|------------|------------|
| 0    | 01-01-1996 | 5.55   | 5.60   | 5.53   | 5.58   | 3.323907   | 985500.0   |
| 1    | 02-01-1996 | 5.47   | 5.57   | 5.29   | 5.37   | 3.198226   | 7470000.0  |
| 2    | 03-01-1996 | 5.13   | 5.25   | 5.10   | 5.20   | 3.095698   | 15160500.0 |
| 3    | 04-01-1996 | 5.20   | 5.33   | 5.14   | 5.30   | 3.153908   | 12397500.0 |
| 4    | 05-01-1996 | 5.30   | 5.28   | 5.19   | 5.20   | 3.097020   | 5008500.0  |
| ...  | ...        | ...    | ...    | ...    | ...    | ...        | ...        |
| 6885 | 10-05-2023 | 423.95 | 426.35 | 422.20 | 425.35 | 425.350006 | 9754485.0  |
| 6886 | 11-05-2023 | 427.25 | 427.40 | 418.90 | 420.40 | 420.399994 | 15982423.0 |
| 6887 | 12-05-2023 | 420.40 | 422.70 | 416.05 | 420.45 | 420.450012 | 11135894.0 |
| 6888 | 15-05-2023 | 420.00 | 428.70 | 419.05 | 427.80 | 427.799988 | 9151102.0  |
| 6889 | 16-05-2023 | 430.00 | 430.00 | 423.50 | 423.95 | 423.950012 | 7994197.0  |

6890 rows × 7 columns

```
In [ ]: columns_to_drop = ['Adj Close']
df.drop(columns=columns_to_drop, inplace = True) # inplace=true implies that r

columns_order = ['Date', 'Open', 'High', 'Low', 'Volume', 'Close']
df = df[columns_order]
df.head(2)
```

```
Out[ ]:
```

|   | Date       | Open | High | Low  | Volume    | Close |
|---|------------|------|------|------|-----------|-------|
| 0 | 01-01-1996 | 5.55 | 5.60 | 5.53 | 985500.0  | 5.58  |
| 1 | 02-01-1996 | 5.47 | 5.57 | 5.29 | 7470000.0 | 5.37  |

```
In [ ]: df['Date'] = pd.to_datetime(df['Date'], format="%d-%m-%Y", errors='coerce')
df['Date'] = df['Date'].dt.strftime('%d-%m-%Y')
df.head(2)
```

```
Out[ ]:
```

|   | Date       | Open | High | Low  | Volume    | Close |
|---|------------|------|------|------|-----------|-------|
| 0 | 01-01-1996 | 5.55 | 5.60 | 5.53 | 985500.0  | 5.58  |
| 1 | 02-01-1996 | 5.47 | 5.57 | 5.29 | 7470000.0 | 5.37  |

## Handling Missing and Duplicate Values

```
In [ ]: df.isnull().sum()
```

```
Out[ ]:
```

|        |    |
|--------|----|
|        | 0  |
| Date   | 0  |
| Open   | 11 |
| High   | 11 |
| Low    | 11 |
| Volume | 11 |
| Close  | 11 |

**dtype:** int64

```
In [ ]: df['Open'] = df['Open'].fillna(df['Open'].mean()) # Replace NaN with the colu
df['High'] = df['High'].fillna(df['High'].mean())
df['Low'] = df['Low'].fillna(df['Low'].mean())
df['Volume'] = df['Volume'].fillna(df['Volume'].mean())
df['Close'] = df['Close'].fillna(df['Close'].mean())
```

```
In [ ]: df.isnull().sum()
```

```
Out[ ]:      0
      Date  0
      Open  0
      High  0
      Low   0
      Volume 0
      Close 0
```

**dtype:** int64

```
In [ ]: # Finding the duplicate values


print(df.duplicated().sum()) # gives the duplicate values in the dataset

0
```

## Data Visualization

```
In [ ]: df.set_index('Date', inplace=True)

# Plot the 'Open' prices with the Date as the x-axis
df['Open'].plot(figsize=(16, 6), title='Open Prices Over Time', legend=True)
plt.xlabel('Date') # Label the x-axis
plt.ylabel('Open Price') # Label the y-axis
plt.grid(True)
plt.show()
```



```
In [ ]: corr = df.corr()['Close']
corr
```


Out[ ]:

|        | Close     |
|--------|-----------|
| Open   | 0.999768  |
| High   | 0.999901  |
| Low    | 0.999885  |
| Volume | -0.306688 |
| Close  | 1.000000  |

**dtype:** float64

```
In [ ]: import seaborn as sns
import matplotlib.pyplot as plt

# Plot the heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(df.corr()[['Close']].sort_values(by='Close', ascending=False),
            annot=True, cmap='coolwarm', fmt='.6f')
plt.title("Correlation of Features with Target Variable")
plt.show()
```




```
In [ ]: # Define colors based on a condition: e.g., 'Close' > 'Open'
        colors = np.where(df['Close'] > df['Open'], 'green', 'red')

        # Scatter plot
        plt.figure(figsize=(8, 6))
        plt.scatter(df['Close'], df['Open'], c=colors, alpha=0.7, s=100)
        plt.title("Correlation between Close and Open Prices")
        plt.xlabel("Close Prices")
        plt.ylabel("Open Prices")

        # Add legend manually
        plt.scatter([], [], c='green', label='Close > Open')
        plt.scatter([], [], c='red', label='Close <= Open')
        plt.legend()

        plt.grid(True)
        plt.show()
```

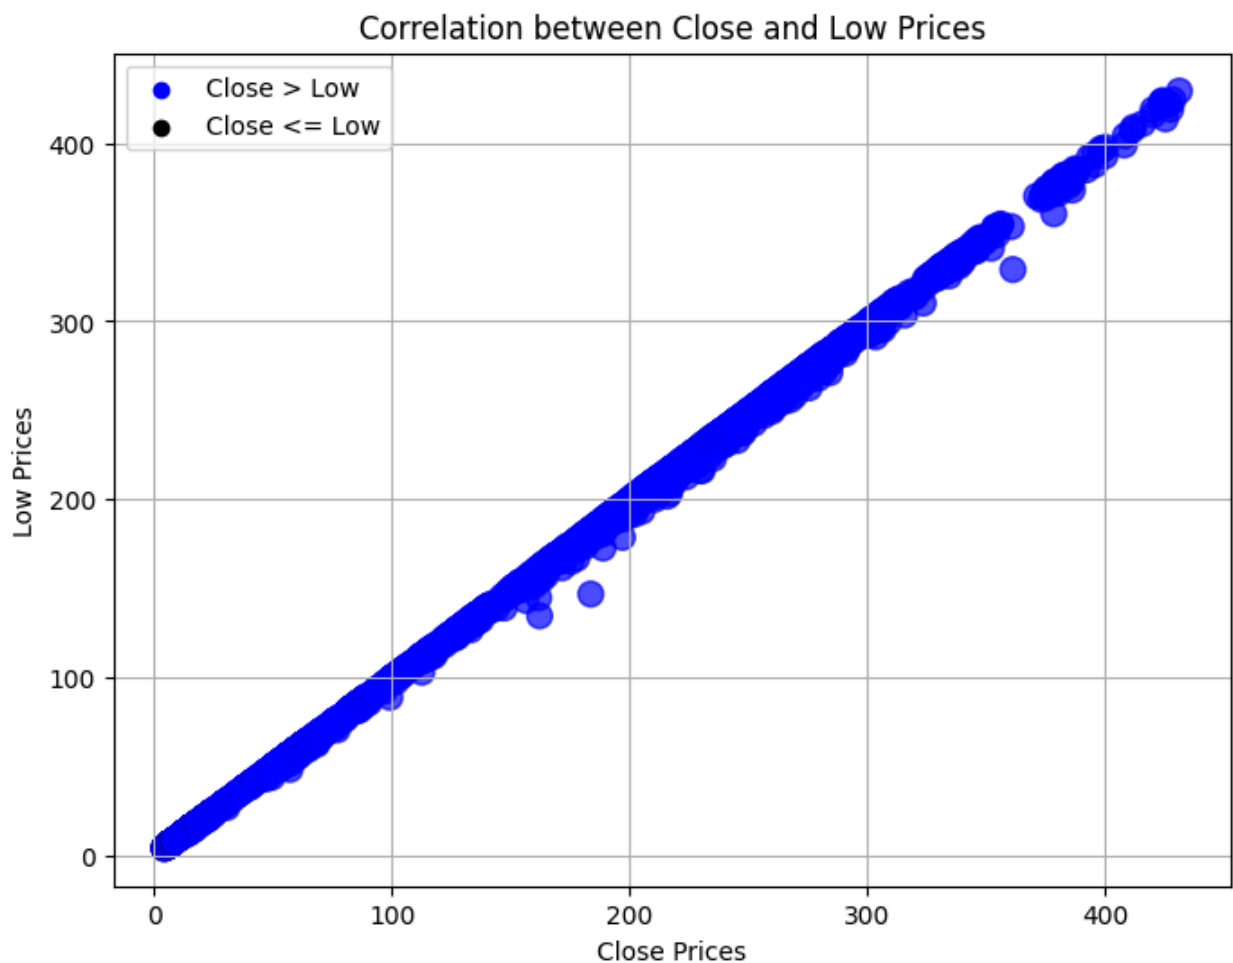


```
In [ ]: # Define colors based on a condition: e.g., 'Close' > 'Low'
        colors = np.where(df['Close'] > df['Low'], 'blue', 'black')

        # Scatter plot
        plt.figure(figsize=(8, 6))
        plt.scatter(df['Close'], df['Low'], c=colors, alpha=0.7, s=100)
        plt.title("Correlation between Close and Low Prices")
        plt.xlabel("Close Prices")
        plt.ylabel("Low Prices")

        # Add legend manually
        plt.scatter([], [], c='blue', label='Close > Low')
        plt.scatter([], [], c='black', label='Close <= Low')
        plt.legend()

        plt.grid(True)
        plt.show()
```



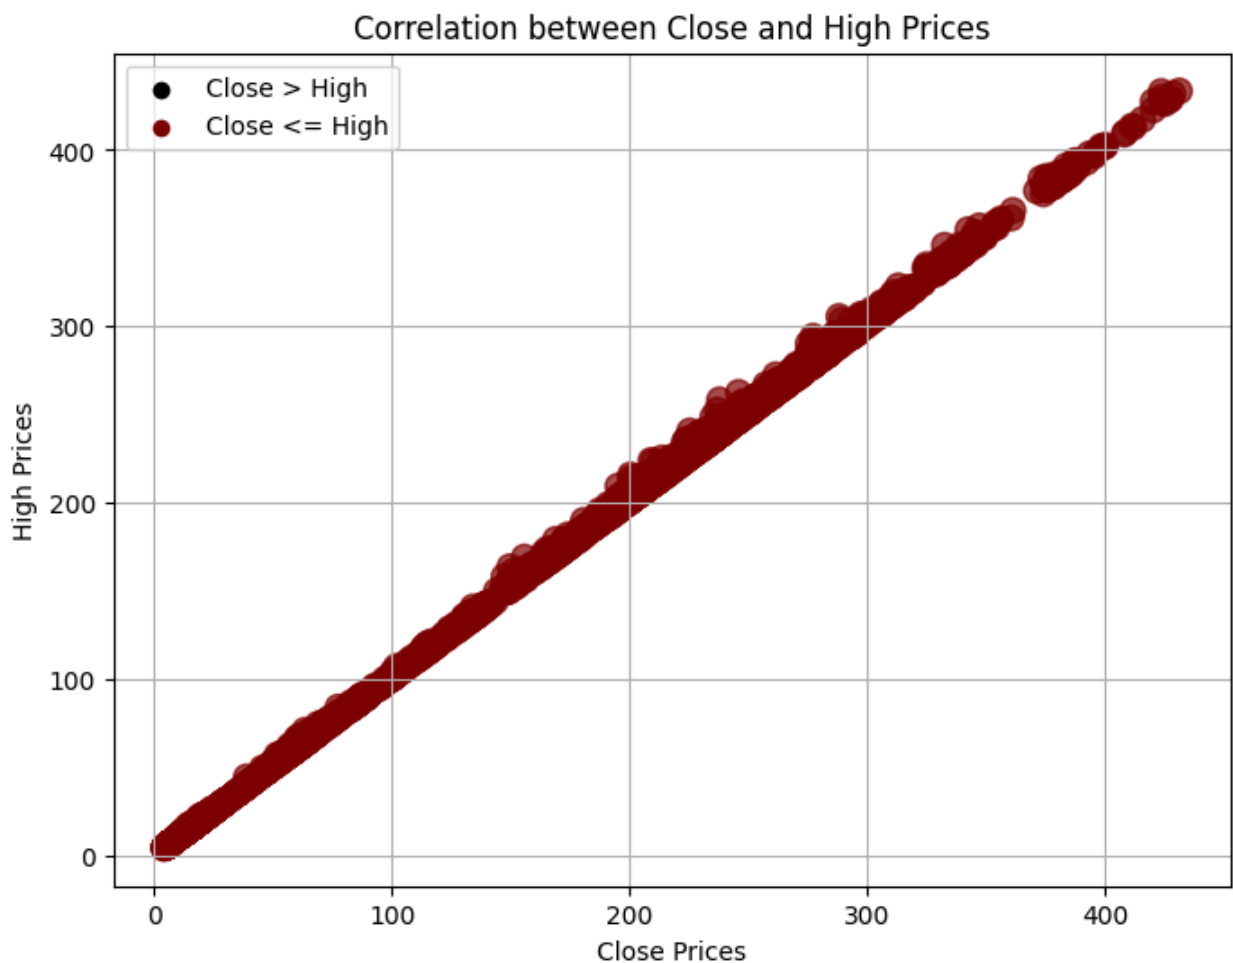
```
In [ ]: # Define colors based on a condition: e.g., 'Close' > 'High'
        colors = np.where(df['Close'] > df['High'], 'black', 'maroon')

        # Scatter plot
        plt.figure(figsize=(8, 6))
        plt.scatter(df['Close'], df['High'], c=colors, alpha=0.7, s=100)
        plt.title("Correlation between Close and High Prices")
        plt.xlabel("Close Prices")
        plt.ylabel("High Prices")

        # Add legend manually
        plt.scatter([], [], c='black', label='Close > High')
        plt.scatter([], [], c='maroon', label='Close <= High')
        plt.legend()

        plt.grid(True)
        plt.show()
```



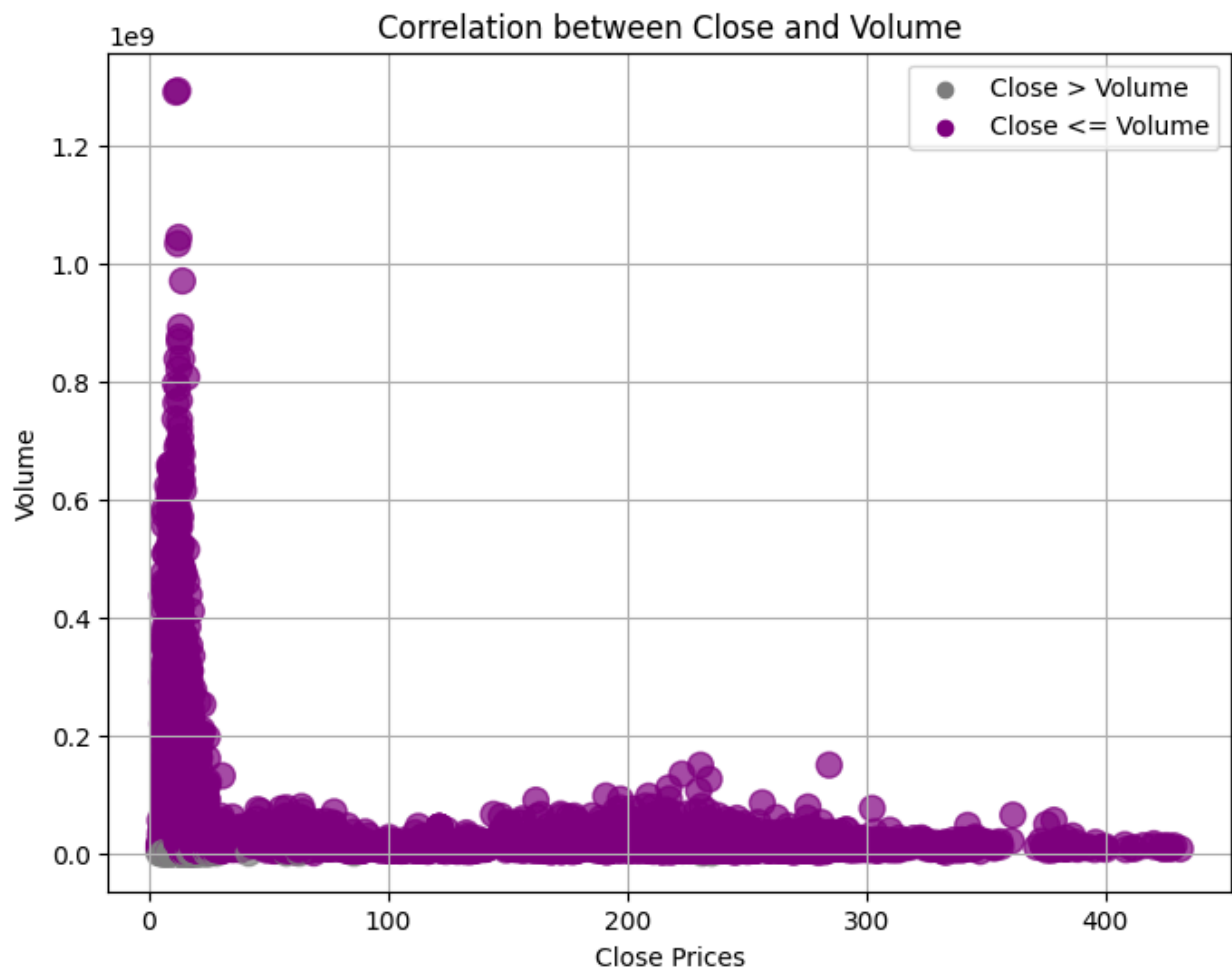


```
In [ ]: # Define colors based on a condition: e.g., 'Close' > 'Volume'
        colors = np.where(df['Close'] > df['Volume'], 'grey', 'purple')

        # Scatter plot
        plt.figure(figsize=(8, 6))
        plt.scatter(df['Close'], df['Volume'], c=colors, alpha=0.7, s=100)
        plt.title("Correlation between Close and Volume")
        plt.xlabel("Close Prices")
        plt.ylabel("Volume")

        # Add legend manually
        plt.scatter([], [], c='grey', label='Close > Volume')
        plt.scatter([], [], c='purple', label='Close <= Volume')
        plt.legend()

        plt.grid(True)
        plt.show()
```



## Splitting the data to training data and test data

```
In [ ]: x = df.drop(columns='Close',axis=1)
        y = df['Close']
```

```
In [ ]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=
```

```
In [ ]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[ ]: ((5512, 4), (1378, 4), (5512,), (1378,))
```

## Data Standardization

```
In [ ]: scaler = StandardScaler()
        scaler.fit(x_train)
```

Out[ ]:

▼ StandardScaler ⓘ ?

StandardScaler()

```
In [ ]: x_train = scaler.transform(x_train)
        x_test = scaler.transform(x_test)
```

```
In [ ]: print(x_train)
```

```
[[ -0.95465118 -0.95771878 -0.95432392 -0.19161583]
 [  1.75622969  1.75417145  1.77971363 -0.37312791]
 [  0.09917298  0.09142023  0.105226   -0.32805496]
 ...
 [  1.22048248  1.20044884  1.22572852 -0.39016636]
 [  1.5024547   1.47614588  1.44361984 -0.35539657]
 [-0.95822282 -0.95147834 -0.95584299  0.27844565]]
```

```
In [ ]: print(x_test)
```

```
[[ -0.99863884 -0.99870078 -0.99961113  0.10617874]
 [  0.58134546  0.57267922  0.57015575 -0.1978014 ]
 [-0.99393931 -0.99553399 -0.99068661 -0.33398784]
 ...
 [-0.06935242 -0.06747813 -0.0782964  -0.36601289]
 [  1.39530525  1.42026135  1.39377542 -0.20869583]
 [-0.54532151 -0.54873712 -0.54256155 -0.3714666  ]]
```

## Model Training using Linear Regression

```
In [ ]: LRmodel = LinearRegression()
        LRmodel.fit(x_train,y_train)
        y_pred = LRmodel.predict(x_test)
        print(y_pred)
```

```
[ 15.79830509 182.28063245  16.48930378 ... 114.17185558 272.4732109
 63.99756266]
```

```
In [ ]: y_pred.shape
```


Out[ ]: (1378,)

```
In [ ]: df = pd.DataFrame({'Actual':y_test,'Predicted':y_pred})
        df.head(10)
```

Out[ ]:                   **Actual**   **Predicted**

| Date       |        |            |
|------------|--------|------------|
| 01-06-2000 | 15.46  | 15.798305  |
| 17-06-2020 | 181.20 | 182.280632 |
| 23-01-2002 | 16.50  | 16.489304  |
| 16-08-2006 | 60.30  | 60.552873  |
| 13-05-2022 | 258.60 | 259.193598 |
| 30-05-2006 | 58.83  | 59.379062  |
| 20-03-2003 | 13.91  | 13.938119  |
| 29-11-2017 | 258.05 | 257.660222 |
| 05-06-2009 | 63.78  | 64.356517  |
| 07-04-2008 | 70.12  | 69.780115  |

```
In [ ]: graph = dframe.head(25)
graph.plot(kind='bar',figsize=(16,5))
plt.title("ITC STOCK : Actual Price vs Predicted Price", fontsize=15)
plt.xlabel("Date", fontsize=12)
plt.ylabel("Price", fontsize=12)
plt.show()
```



## Model Evaluation for Linear Regression Model

```
In [ ]: r2_score(y_test, y_pred)
```

Out[ ]: 0.999884066238908

```
In [ ]: print('Mean Absolute Error',metrics.mean_absolute_error(y_test,y_pred))
        print('Mean Squared Error',metrics.mean_squared_error(y_test,y_pred))
        import math
        print ('Root Mean Squared Error',math.sqrt(metrics.mean_squared_error(y_test,y
```

Mean Absolute Error 0.564774466912428  
Mean Squared Error 1.3041262317038838  
Root Mean Squared Error 1.1419834638487039

## Predictive Model using Linear Regression

```
In [ ]: input_data = (120,130,123,980000)

        #changing input data to a numpy array
        input_data_as_numpy_array = np.asarray(input_data)

        # reshape the numpy array
        input_data_resaped = input_data_as_numpy_array.reshape(1,-1)

        # standardizing the data
        std_data = scaler.transform(input_data_resaped)

        prediction = LRmodel.predict(std_data)
        print(prediction)
        # predicting the close price
```

[130.12984327]

```
In [ ]: input_data = (465.45,466.45,462.60,8823805)

        #changing input data to a numpy array
        input_data_as_numpy_array = np.asarray(input_data)

        # reshape the numpy array
        input_data_resaped = input_data_as_numpy_array.reshape(1,-1)

        # standardizing the data
        std_data = scaler.transform(input_data_resaped)

        prediction = LRmodel.predict(std_data)
        print(prediction)
        # predicting the close price
```

[463.87728268]

## Model Training using Support Vector Regressor

```
In [ ]: svr_model = SVR(kernel='linear', C=50, epsilon=0.2)
        svr_model.fit(x_train, y_train)
```

Out[ ]:

SVR

SVR(C=50, epsilon=0.2, kernel='linear')

```
In [ ]: prediction = svr_model.predict(x_test)
print(prediction)
```


```
[ 15.78699691 182.18550435 16.50436631 ... 114.10188452 272.44207497
 64.00406944]
```

```
In [ ]: dframeS = pd.DataFrame({'Actual':y_test,'Predicted':prediction})
dframeS.head(10)
```

Out[ ]:

|            | Actual | Predicted  |
|------------|--------|------------|
| Date       |        |            |
| 01-06-2000 | 15.46  | 15.786997  |
| 17-06-2020 | 181.20 | 182.185504 |
| 23-01-2002 | 16.50  | 16.504366  |
| 16-08-2006 | 60.30  | 60.593271  |
| 13-05-2022 | 258.60 | 259.317527 |
| 30-05-2006 | 58.83  | 59.337018  |
| 20-03-2003 | 13.91  | 13.951485  |
| 29-11-2017 | 258.05 | 257.735181 |
| 05-06-2009 | 63.78  | 64.132758  |
| 07-04-2008 | 70.12  | 69.751918  |

```
In [ ]: graphS = dframeS.head(25)
graphS.plot(kind='bar',figsize=(16,5))
plt.title("ITC STOCK : Actual Price vs Predicted Price", fontsize=15)
plt.xlabel("Date", fontsize=12)
plt.ylabel("Price", fontsize=12)
plt.show()
```



## Model Evaluation for Support Vector Regressor Model

```
In [ ]: r2_score(y_test, prediction)
```

```
Out[ ]: 0.9998759043005752
```

```
In [ ]: print('Mean Absolute Error',metrics.mean_absolute_error(y_test,prediction))
print('Mean Squared Error',metrics.mean_squared_error(y_test,prediction))
import math
print ('Root Mean Squared Error',math.sqrt(metrics.mean_squared_error(y_test,p
```

Mean Absolute Error 0.5604446370532994

Mean Squared Error 1.3959389856524873

Root Mean Squared Error 1.181498618557164

## Predictive Model for Support Vector Regressor

```
In [ ]: input_data = (120,130,123,980000)

#changing input data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the numpy array
input_data_resaped = input_data_as_numpy_array.reshape(1,-1)

# standardizing the data
std_data = scaler.transform(input_data_resaped)

pred = svr_model.predict(std_data)
print(pred)
# predicting the close price
```

[130.41050458]

```
In [ ]: input_data = (465.45,466.45,462.60,8823805)

#changing input data to a numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the numpy array
input_data_resaped = input_data_as_numpy_array.reshape(1,-1)

# standardizing the data
std_data = scaler.transform(input_data_resaped)

pred = svr_model.predict(std_data)
print(pred)
# predicting the close price
```

[463.93927574]